

COP4610

Introduction to Operating Systems


Project #3:

Implementing a FAT32
File System

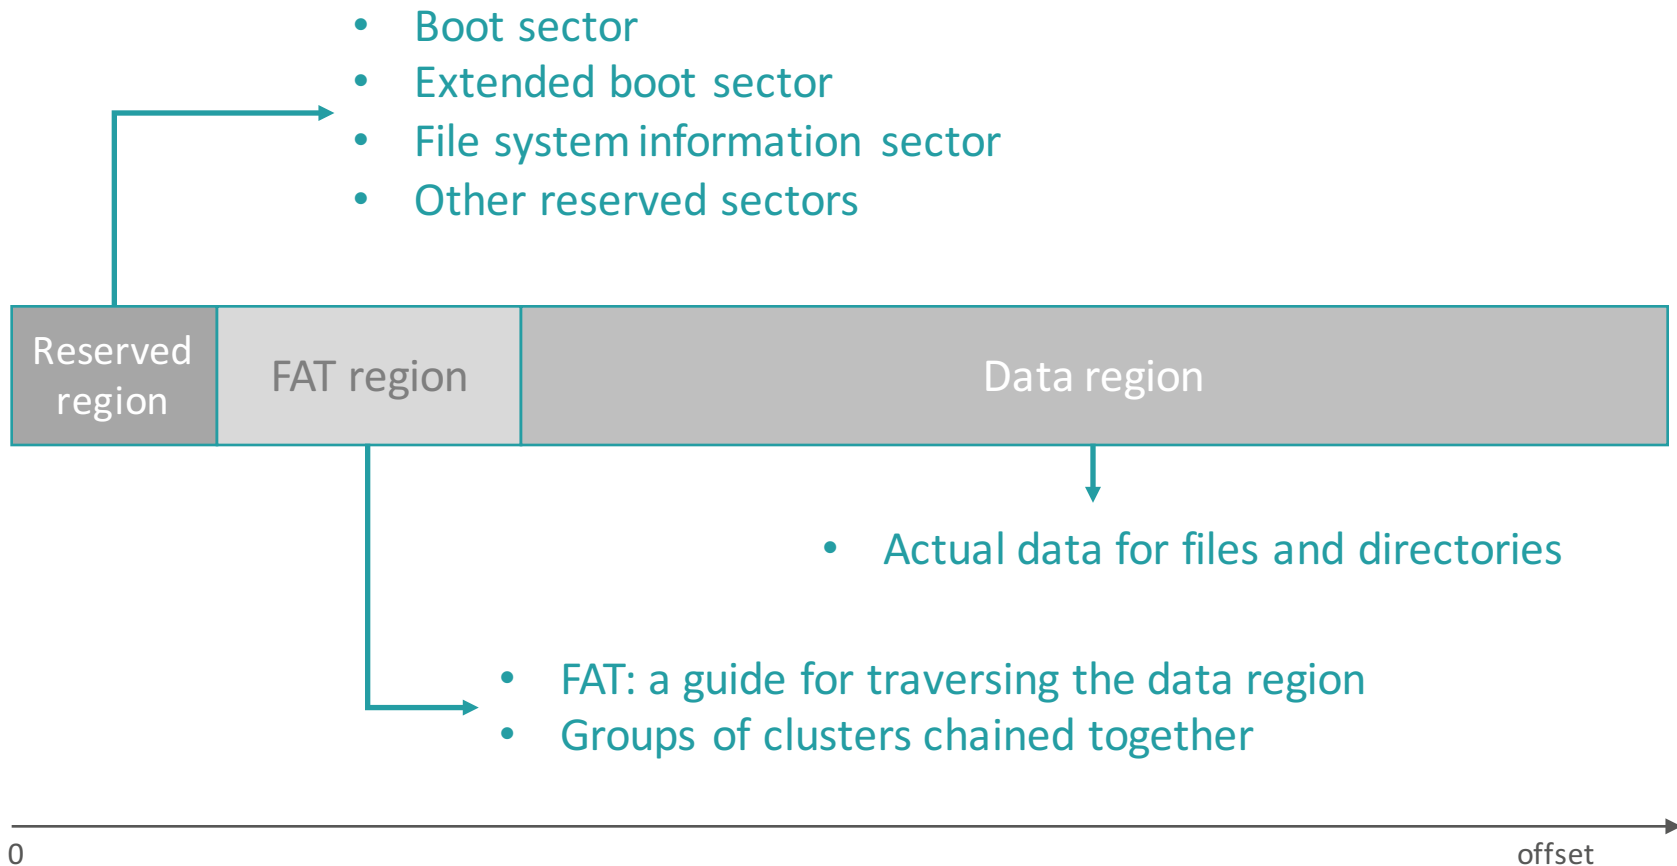
Outline

- Review
- C functions, data types and program structure
- Reserved region: BPB
- Accessing data
 - File contents
 - Directories
- Program commands
 - `exit`
 - `info`
 - `ls DIRNAME`

Review



Review



Useful C functions

- `int open(const char *path, int oflag, ...);`
- `lseek(int fildes, off_t offset, int whence);`
- `ssize_t read(int fildes, void *buf, size_t nbyte);`
- `ssize_t write(int fildes, const void *buf, size_t nbyte);`
- `int close(int fildes);`

Useful C data types

- `unsigned char`
 - 8 bits without sign
- `unsigned short`
 - 16 bits without sign
- `unsigned int`
 - 32 bits without sign
- Be sure to use **unsigned** types when working with the volume!

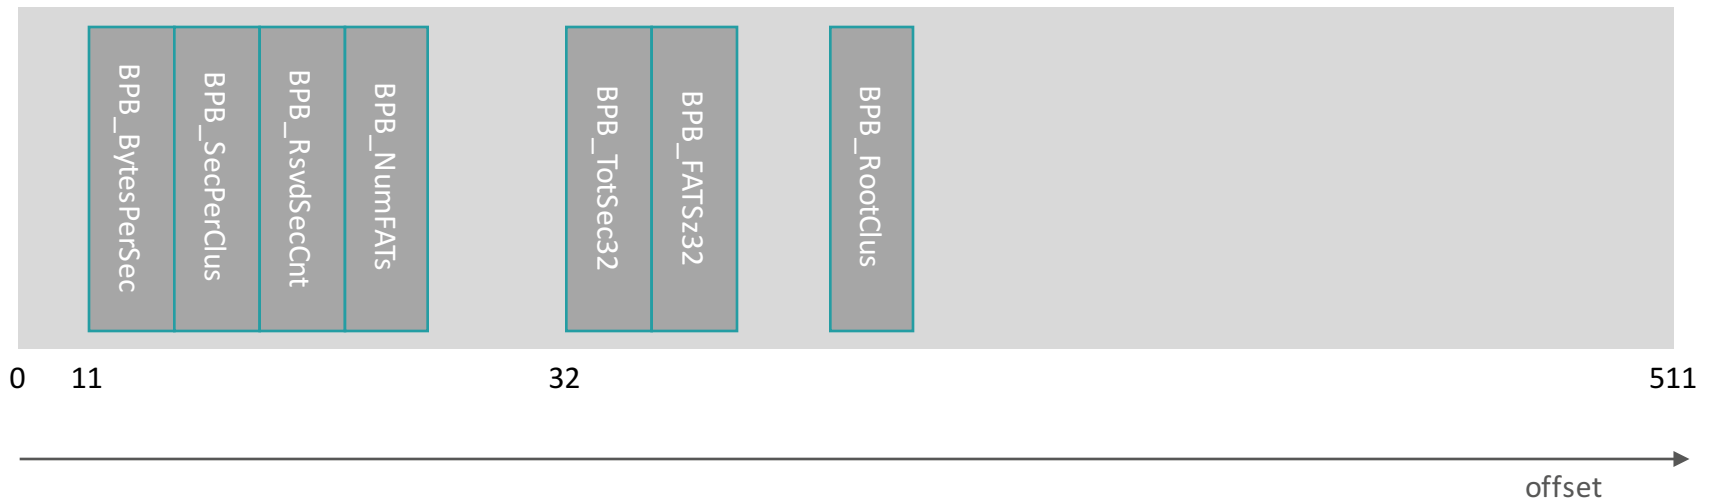
Program structure

```
while (command != "exit") {  
    print("$ ");  
    read_from_stdin(command);  
    if (command == "info") {  
        do_something();  
    }  
    else if (command == "ls") {  
        do_something2();  
    }  
    ...  
}
```

BPB (BIOS Parameter Block)

- Located at the first sector of the volume
- Part of the Reserved Region (first sector(s))
- Contains important information about the file system
 - Bytes per sector (BPB_BytsPerSec)
 - Sectors per cluster (BPB_SecPerClus)
 - Reserved region size (BPB_RsvdSecCnt)
 - Number of FATs (BPB_NumFATs)
 - FAT size (BPB_FATSz32)
 - Root cluster (BPB_RootClus)
 - Total sectors (BPB_TotSec32)

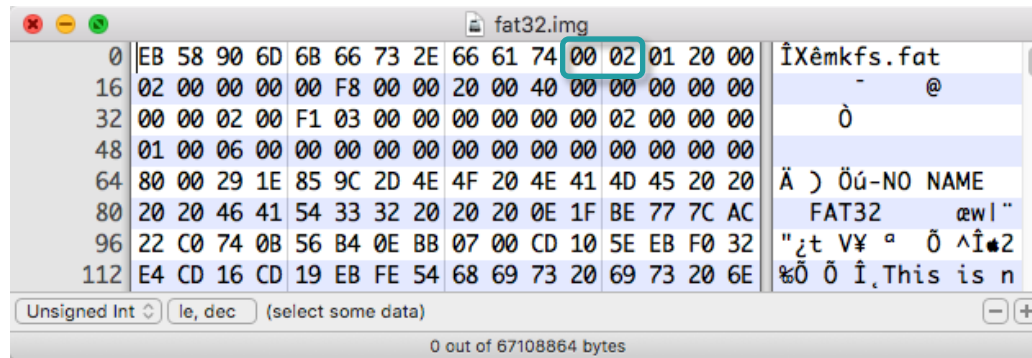
BPB (BIOS Parameter Block)



- Refer to FAT32 specification for details

BPB (BIOS Parameter Block)

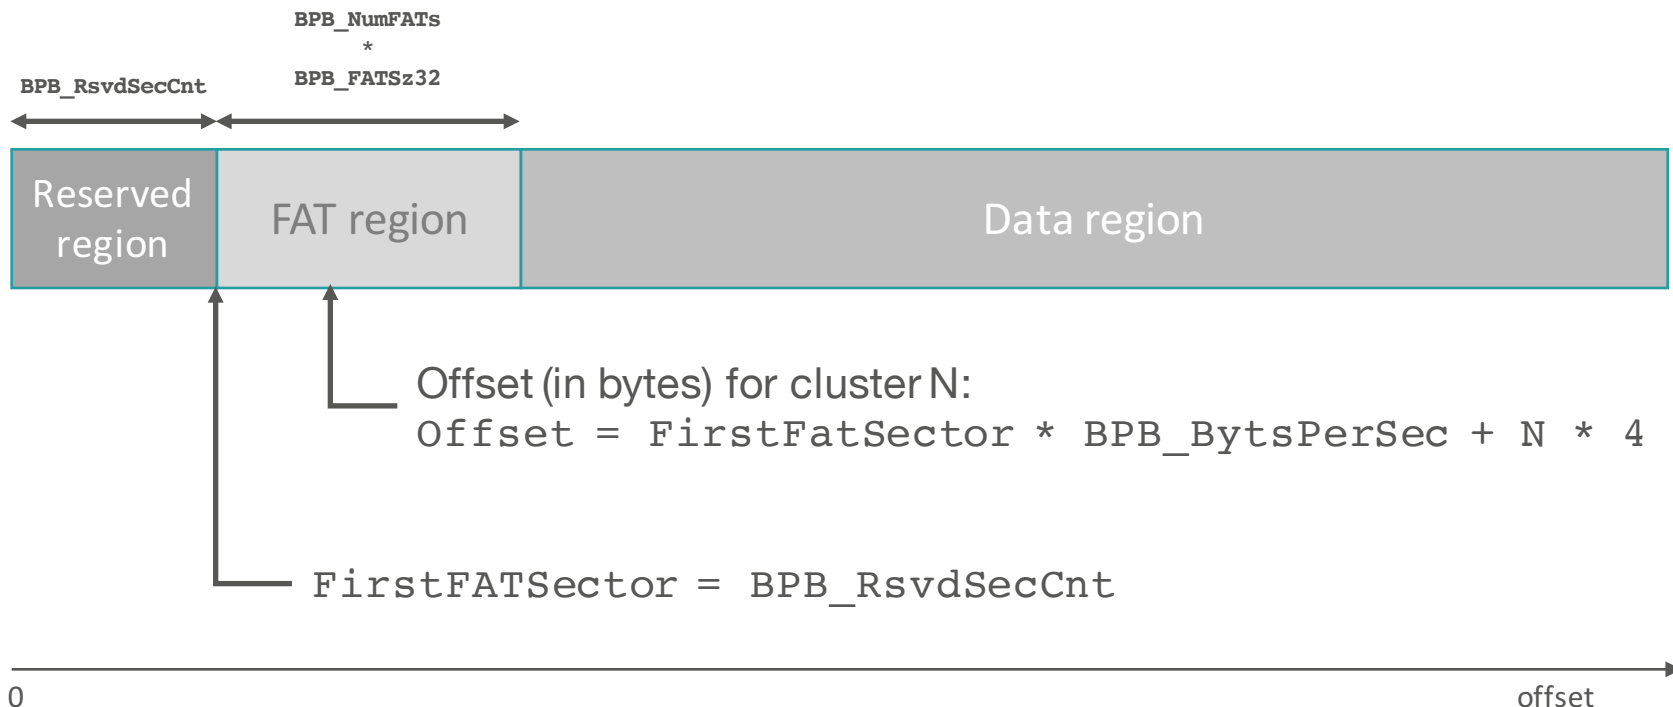
- Example: Bytes per sector
 - Field in FAT32 spec: BPB_BytsPerSec
 - Offset: 11 bytes
 - Size: 2 bytes



- Remember: Little endian
 - 00 02 → 0x0200 = 512 bytes per sector

Traversing the FAT

- Find FAT's offset, then find offset within FAT for a given cluster

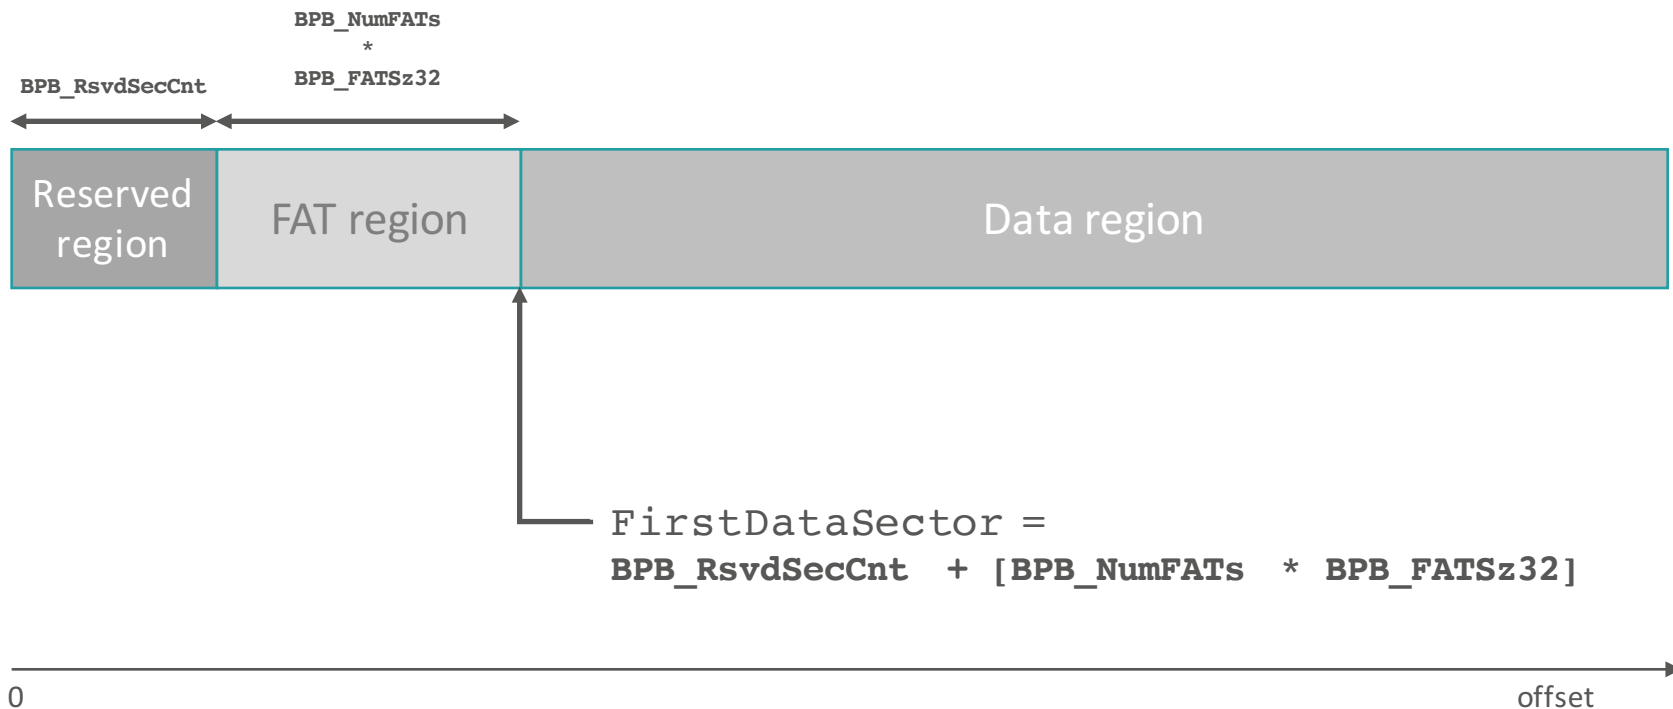


Traversing the FAT

- Get the offset of the FAT (in bytes)
- For a given cluster, find its entry
 - Find the cluster's offset
 - Read the entry (4 bytes) at that offset
 - If the entry is 0x0FFFFFFF8 to 0x0FFFFFFFE (or 0xFFFFFFFFF), it means it is the last cluster. Stop
 - If not, the read entry is the next cluster. Repeat.

Accessing File Contents

- Find the first data sector



Accessing File Contents

- Then find the offset for a given cluster N (within the data region)
 - Intuition: ~~$N * \text{BPB_SecPerClus}$~~
 - But the first cluster was 2!
 - And the data begins at the beginning of the data region
 - $\rightarrow (N - 2) * \text{BPB_SecPerClus}$
- Now combine with the first data region
 - We get the offset (sector) of cluster N within the volume, given its number:

$$\text{Offset} = \text{FirstDataSector} + (N - 2) * \text{BPB_SecPerClus}$$

Accessing File Contents

- Steps
 - Get the offset for the first cluster (previous slide)
 - Calculate the byte offset
 - $\text{ByteOffset} = \text{Offset} * \text{BPB_BytsPerSec}$
 - Read/write data at that offset within the image file
 - Traverse FAT when needed (i.e.: you reached the end of the cluster and need to read more data)
 - Don't go beyond the file size
 - Don't go beyond cluster size without traversing
 - Clusters will (most likely) not be one after each other

Accessing File Contents

Root cluster
↓

0000000	XXXXXXXX	XXXXXXXX	00000009	00000004
0000004	00000005	00000007	00000000	00000008
0000008	FFFFFFFF	0000000A	0000000B	00000011
000000C	0000000D	0000000E	FFFFFFFF	00000010
0000010	00000012	FFFFFFFF	00000013	00000014
0000014	00000015	00000016	FFFFFFFF	00000000
0000018	00000000	00000000	00000000	00000000
000001C	00000000	00000000	00000000	00000000
0000020	00000000	00000000	00000000	00000000
...

```
while not_last_cluster():  
    read/write_current();  
    get_next_cluster();  
end while;
```

FAT entry indicates where the next cluster is located

- If 0x0: cluster not allocated
- If 0xFFFFFFFF (or 0x0FFFFFFF8 to 0x0FFFFFFFE): last cluster
- Otherwise: indicates next cluster of the file/directory

Accessing Directories

- Directories are also considered files
- Content is directory entries (DIRENTRY)
 - Fixed size structure (32B) that contains information about objects in the directory (files, subdirectories)
 - See section 6 of FAT32 spec
- How to get info about the directory entries
 - Find directory's first cluster (and its offset)
 - Use a structure to read the data at the offset
 - Does `read()` sound familiar?
 - Traverse the FAT when needed

Accessing Directories

DIRENTRY

```
struct DIRENTRY {  
    unsigned char DIR_name[11];  
    unsigned char DIR_Attributes;  
    ...  
} __attribute__((packed));
```

- A `struct DIRENTRY*` can be used as parameter for `read()` and `write()`
- Be sure to use `__attribute__((packed))` in order to avoid attribute padding

Accessing Directories

- Not all read directory entries are actual directory entries
 - Some are empty entries
 - Check first byte (`DIR_name[0]`)
 - 0x0: empty entry and no other entries after it
 - 0xE5: empty entry but more entries after it
 - Otherwise: entry is taken
 - Some of them are long DIR name entries
 - Can be ignored
 - Use `ATTR_LONG_NAME` mask on `DIR_Attributes` to determine whether it is a `DIRENTRY` or long DIR name
 - Check FAT32 specs

Accessing Directories

- DIR entries

	001003F0	00 00 00 00	00 00 00 00	00 00 00 00	00 00 00 00
	00100400	41 6C 00 6F	00 6E 00 67	00 66 00 0F	00 97 69 00	Al.o.n.g.f....i.
	00100410	6C 00 65 00	00 00 FF FF	FF FF 00 00	FF FF FF FF	l.e.....
	00100420	4C 4F 4E 47	46 49 4C 45	20 20 20 20	00 64 04 8E	LONGFILE .d..
	00100430	78 4E 49 4F	00 00 04 8E	78 4E 03 00	76 5B 03 00	xNIO....xN..v[..
Long DIR Name Entries (ignore)	00100440	41 68 00 65	00 6C 00 6C	00 6F 00 0F	00 14 00 00	Ah.e.l.l.o.....
	00100450	FF FF FF FF	FF FF FF FF	FF FF 00 00	FF FF FF FF
	00100460	48 45 4C 4C	4F 20 20 20	20 20 20 20	00 64 04 8E	HELLO .d..
	00100470	78 4E 49 4F	00 00 04 8E	78 4E B1 01	0A 00 00 00	xNIO....xN.....
	00100480	41 62 00 6C	00 75 00 65	00 00 00 0F	00 55 FF FF	Ab.l.u.e.....U..
	00100490	FF FF FF FF	FF FF FF FF	FF FF 00 00	FF FF FF FF
	001004A0	42 4C 55 45	20 20 20 20	20 20 20 10	00 64 04 8E	BLUE .d..
	001004B0	78 4E 78 4E	00 00 04 8E	78 4E B2 01	00 00 00 00	xNxN....xN.....
	001004C0	41 67 00 72	00 65 00 65	00 6E 00 0F	00 42 00 00	Ag.r.e.e.n...B..
	001004D0	FF FF FF FF	FF FF FF FF	FF FF 00 00	FF FF FF FF
	001004E0	47 52 45 45	4E 20 20 20	20 20 20 10	00 64 04 8E	GREEN .d..
	001004F0	78 4E 78 4E	00 00 04 8E	78 4E B3 01	00 00 00 00	xNxN....xN.....
	00100500	41 72 00 65	00 64 00 00	00 FF FF 0F	00 37 FF FF	Ar.e.d.....7..
	00100510	FF FF FF FF	FF FF FF FF	FF FF 00 00	FF FF FF FF
	00100520	52 45 44 20	20 20 20 20	20 20 20 10	00 00 05 8E	RED
	00100530	78 4E 78 4E	00 00 05 8E	78 4E B4 01	00 00 00 00	xNxN....xN.....
	00100540	00 00 00 00	00 00 00 00	00 00 00 00	00 00 00 00
	00100550	00 00 00 00	00 00 00 00	00 00 00 00	00 00 00 00
	--- fat32.img --0x100550/0x4000000-----					

DIR entry for HELLO file

DIR entry for GREEN directory

Program Commands

exit

- Releases all allocated resources
- Safely exits the program

Program Commands

info

- Prints important data about the file system
- Read from BPB
- Two options
 - Move to the offset of the volume and read values individually
 - Create a structure that represents the BPB and read by passing the whole structure

Program Commands

info

- Be careful!
 - If you get the wrong data from the BPB you might not be able to continue until you fix it
 - All the other commands depend on this data
- Do the values make sense?
 - verify by using Hexedit and the FAT32 spec

Program Commands

ls DIRNAME

- Show the names of the files/directories under the directory given by DIRNAME (local)
- Read the current directory's entries until you find DIRNAME's (and its first cluster!)
 - Then read DIRNAME's entries and print their names
- If no argument given, print the names of the entries at the current directory
 - Read entries for current directory's cluster

Program Commands

ls DIRNAME

- Hint: Once current directory's entries were read, find DIRNAME's cluster within current dir
 - And repeat the process
- If DIRNAME does not exist, print an error

Questions?